

Lab 3 : Binary Classification of Heart Disease of Patients using Deep Neural Network

Dinesh Kumar K

225229108

1. Load the dataset

```
In [1]: import pandas as pd
```

```
In [2]: df = pd.read_csv("heart_data.csv")
```

```
In [3]: df.head()
```

Out[3]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1

```
In [4]: df.shape
```

Out[4]: (303, 14)

```
In [5]: df.size
```

Out[5]: 4242

```
In [6]: df.columns
```

Out[6]: Index(['age', 'sex', 'cp', 'trestbps', 'chol', 'fbs', 'restecg', 'thalach', 'exang', 'oldpeak', 'slope', 'ca', 'thal', 'target'], dtype='object')

2. Split the dataset

```
In [9]: X = df.drop('target', axis=1)
        y = df['target']
```

```
In [10]: from sklearn.model_selection import train_test_split
         X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_s
```

```
In [11]: X_train.shape
```

```
Out[11]: (242, 13)
```

```
In [12]: X_test.shape
```

```
Out[12]: (61, 13)
```

3. Create a neural network

```
In [13]: #Input size = No. of features in X_train = 13  
#No. of neurons/units in the Dense layer = 8, with Relu activation function  
#No. of neurons/units in output layer = 1, with sigmoid activation function
```

```
In [14]: from tensorflow.keras.models import Sequential  
         from tensorflow.keras.layers import Dense
```

```
In [15]: model = Sequential()  
         model.add(Dense(8, input_dim=13, activation='relu'))  
         model.add(Dense(1, activation='sigmoid'))
```

4. Compile your model

```
In [16]: from tensorflow import keras
```

```
In [17]: optimizer = keras.optimizers.RMSprop(learning_rate=0.001)
```

```
In [18]: model.compile(loss='mse', optimizer=optimizer, metrics=['accuracy'])  
         model.fit(X_train, y_train, epochs=10, batch_size=30, verbose=1)
```

```
Epoch 1/10  
9/9 [=====] - 2s 12ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 2/10  
9/9 [=====] - 0s 1ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 3/10  
9/9 [=====] - 0s 1ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 4/10  
9/9 [=====] - 0s 2ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 5/10  
9/9 [=====] - 0s 2ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 6/10  
9/9 [=====] - 0s 1ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 7/10  
9/9 [=====] - 0s 1ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 8/10  
9/9 [=====] - 0s 1ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 9/10  
9/9 [=====] - 0s 2ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 10/10  
9/9 [=====] - 0s 2ms/step - loss: 0.4504 - accuracy:  
0.5496
```

```
Out[18]: <keras.callbacks.History at 0x28abb93efb0>
```

```
In [19]: model.evaluate(X_test, y_test)
```

```
2/2 [=====] - 0s 3ms/step - loss: 0.4754 - accuracy:  
0.5246
```

```
Out[19]: [0.4754098355770111, 0.5245901346206665]
```

5. Print the summary of the model

In [20]: `model.summary()`

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 8)	112
dense_1 (Dense)	(None, 1)	9
Total params: 121		
Trainable params: 121		
Non-trainable params: 0		

6. Train the model

In [21]: `model.compile(loss='mse', optimizer=optimizer, metrics=['accuracy'])`
`model.fit(X_train, y_train, epochs=200, batch_size=10, verbose=1)`

```

y: 0.5496
Epoch 2/200
25/25 [=====] - 0s 1ms/step - loss: 0.4504 - accurac
y: 0.5496
Epoch 3/200
25/25 [=====] - 0s 1ms/step - loss: 0.4504 - accurac
y: 0.5496
Epoch 4/200
25/25 [=====] - 0s 1ms/step - loss: 0.4504 - accurac
y: 0.5496
Epoch 5/200
25/25 [=====] - 0s 1ms/step - loss: 0.4504 - accurac
y: 0.5496
Epoch 6/200
25/25 [=====] - 0s 1ms/step - loss: 0.4504 - accurac
y: 0.5496
Epoch 7/200
25/25 [=====] - 0s 1ms/step - loss: 0.4504 - accurac
y: 0.5496
Epoch 8/200
-- -- -- -- --

```

In [22]: `model.evaluate(X_test, y_test)`

```

2/2 [=====] - 0s 4ms/step - loss: 0.4754 - accuracy:
0.5246

```

Out[22]: [0.4754098355770111, 0.5245901346206665]

7. Save the trained model

In [23]: `history = model.fit(X_train, y_train, validation_split=0.2, epochs=100, batch_size=20)`

```
Epoch 1/100
20/20 [=====] - 0s 6ms/step - loss: 0.4560 - accuracy: 0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
Epoch 2/100
20/20 [=====] - 0s 3ms/step - loss: 0.4560 - accuracy: 0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
Epoch 3/100
20/20 [=====] - 0s 3ms/step - loss: 0.4560 - accuracy: 0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
Epoch 4/100
20/20 [=====] - 0s 3ms/step - loss: 0.4560 - accuracy: 0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
Epoch 5/100
20/20 [=====] - 0s 3ms/step - loss: 0.4560 - accuracy: 0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
Epoch 6/100
20/20 [=====] - 0s 3ms/step - loss: 0.4560 - accuracy: 0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
Epoch 7/100
20/20 [=====] - 0s 3ms/step - loss: 0.4560 - accuracy: 0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
```

8. Evaluate

In [24]: `model.evaluate(X_test, y_test)`

```
2/2 [=====] - 0s 3ms/step - loss: 0.4754 - accuracy: 0.5246
```

Out[24]: `[0.4754098355770111, 0.5245901346206665]`


9. Print the model accuracy


In [28]: `history.history.keys()`

Out[28]: `dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])`

In [29]: `import matplotlib.pyplot as plt`

```
In [30]: plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['Train', 'Validation'])
plt.show()
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['Train', 'Validation'])
plt.show()
```





10. Do further experiments

```
In [31]: model1 = Sequential()  
model1.add(Dense(16, input_dim=13, activation='relu'))  
model1.add(Dense(8, activation='relu'))  
model1.add(Dense(1, activation='sigmoid'))
```

```
In [32]: model1.compile(loss='mse', optimizer=optimizer, metrics=['accuracy'])  
         model1.fit(X_train, y_train, epochs=10, batch_size=30, verbose=1)
```

```
Epoch 1/10  
9/9 [=====] - 1s 2ms/step - loss: 0.5496 - accuracy:  
0.4504  
Epoch 2/10  
9/9 [=====] - 0s 2ms/step - loss: 0.5496 - accuracy:  
0.4504  
Epoch 3/10  
9/9 [=====] - 0s 1ms/step - loss: 0.5496 - accuracy:  
0.4504  
Epoch 4/10  
9/9 [=====] - 0s 1ms/step - loss: 0.5496 - accuracy:  
0.4504  
Epoch 5/10  
9/9 [=====] - 0s 2ms/step - loss: 0.5496 - accuracy:  
0.4504  
Epoch 6/10  
9/9 [=====] - 0s 1ms/step - loss: 0.5496 - accuracy:  
0.4504  
Epoch 7/10  
9/9 [=====] - 0s 2ms/step - loss: 0.5496 - accuracy:  
0.4504  
Epoch 8/10  
9/9 [=====] - 0s 2ms/step - loss: 0.5496 - accuracy:  
0.4504  
Epoch 9/10  
9/9 [=====] - 0s 1ms/step - loss: 0.5496 - accuracy:  
0.4504  
Epoch 10/10  
9/9 [=====] - 0s 2ms/step - loss: 0.5496 - accuracy:  
0.4504
```

```
Out[32]: <keras.callbacks.History at 0x28ac0d19cf0>
```

```
In [33]: model1.evaluate(X_test, y_test)
```

```
2/2 [=====] - 0s 2ms/step - loss: 0.5246 - accuracy:  
0.4754
```

```
Out[33]: [0.5245901346206665, 0.4754098355770111]
```


In [34]: `history1 = model.fit(X_train, y_train, validation_split=0.2, epochs=100, batch_si`

```
Epoch 1/100
7/7 [=====] - 0s 15ms/step - loss: 0.4560 - accurac
y: 0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
Epoch 2/100
7/7 [=====] - 0s 6ms/step - loss: 0.4560 - accuracy:
0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
Epoch 3/100
7/7 [=====] - 0s 6ms/step - loss: 0.4560 - accuracy:
0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
Epoch 4/100
7/7 [=====] - 0s 6ms/step - loss: 0.4560 - accuracy:
0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
Epoch 5/100
7/7 [=====] - 0s 6ms/step - loss: 0.4560 - accuracy:
0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
Epoch 6/100
7/7 [=====] - 0s 5ms/step - loss: 0.4560 - accuracy:
0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
Epoch 7/100
7/7 [=====] - 0s 5ms/step - loss: 0.4560 - accuracy:
0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
```

In [35]: `model1.summary()`

Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 16)	224
dense_3 (Dense)	(None, 8)	136
dense_4 (Dense)	(None, 1)	9
Total params: 369		
Trainable params: 369		
Non-trainable params: 0		

In [36]: `ls = history1.history`

```
In [37]: new = pd.DataFrame.from_dict(ls)
new
```

Out[37]:

	loss	accuracy	val_loss	val_accuracy
0	0.455959	0.544041	0.428571	0.571429
1	0.455959	0.544041	0.428571	0.571429
2	0.455959	0.544041	0.428571	0.571429
3	0.455959	0.544041	0.428571	0.571429
4	0.455959	0.544041	0.428571	0.571429
...
95	0.455959	0.544041	0.428571	0.571429
96	0.455959	0.544041	0.428571	0.571429
97	0.455959	0.544041	0.428571	0.571429
98	0.455959	0.544041	0.428571	0.571429
99	0.455959	0.544041	0.428571	0.571429

100 rows × 4 columns

```
In [38]: model2 = Sequential()
model2.add(Dense(32, input_dim=13, activation='relu'))
model2.add(Dense(16, activation='relu'))
model2.add(Dense(8, activation='relu'))
model2.add(Dense(1, activation='sigmoid'))
```

```
In [39]: model2.compile(loss='mse', optimizer=optimizer, metrics=['accuracy'])  
model2.fit(X_train, y_train, epochs=10, batch_size=30, verbose=1)
```

```
Epoch 1/10  
9/9 [=====] - 1s 2ms/step - loss: 0.5427 - accuracy:  
0.4421  
Epoch 2/10  
9/9 [=====] - 0s 2ms/step - loss: 0.3125 - accuracy:  
0.4876  
Epoch 3/10  
9/9 [=====] - 0s 2ms/step - loss: 0.2681 - accuracy:  
0.5041  
Epoch 4/10  
9/9 [=====] - 0s 2ms/step - loss: 0.2501 - accuracy:  
0.5868  
Epoch 5/10  
9/9 [=====] - 0s 2ms/step - loss: 0.2387 - accuracy:  
0.5950  
Epoch 6/10  
9/9 [=====] - 0s 2ms/step - loss: 0.2403 - accuracy:  
0.6074  
Epoch 7/10  
9/9 [=====] - 0s 2ms/step - loss: 0.2658 - accuracy:  
0.5661  
Epoch 8/10  
9/9 [=====] - 0s 2ms/step - loss: 0.2274 - accuracy:  
0.6240  
Epoch 9/10  
9/9 [=====] - 0s 2ms/step - loss: 0.2452 - accuracy:  
0.6281  
Epoch 10/10  
9/9 [=====] - 0s 2ms/step - loss: 0.2221 - accuracy:  
0.6240
```

```
Out[39]: <keras.callbacks.History at 0x28ac1fbadd0>
```

```
In [40]: model2.evaluate(X_test, y_test)
```

```
2/2 [=====] - 0s 3ms/step - loss: 0.3232 - accuracy:  
0.5082
```

```
Out[40]: [0.3232260048389435, 0.5081967115402222]
```

```
In [41]: model2.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
dense_5 (Dense)	(None, 32)	448
dense_6 (Dense)	(None, 16)	528
dense_7 (Dense)	(None, 8)	136
dense_8 (Dense)	(None, 1)	9
=====		
Total params: 1,121		
Trainable params: 1,121		
Non-trainable params: 0		
=====		

```
In [42]: model3 = Sequential()
model3.add(Dense(64, input_dim=13, activation='relu'))
model3.add(Dense(32, activation='relu'))
model3.add(Dense(16, activation='relu'))
model3.add(Dense(8, activation='relu'))
model3.add(Dense(1, activation='sigmoid'))
```

```
In [43]: model3.compile(loss='mse', optimizer=optimizer, metrics=['accuracy'])
model3.fit(X_train, y_train, epochs=10, batch_size=30, verbose=1)
```

```
Epoch 1/10
9/9 [=====] - 1s 2ms/step - loss: 0.4504 - accuracy: 0.5496
Epoch 2/10
9/9 [=====] - 0s 1ms/step - loss: 0.4504 - accuracy: 0.5496
Epoch 3/10
9/9 [=====] - 0s 2ms/step - loss: 0.4504 - accuracy: 0.5496
Epoch 4/10
9/9 [=====] - 0s 2ms/step - loss: 0.4504 - accuracy: 0.5496
Epoch 5/10
9/9 [=====] - 0s 2ms/step - loss: 0.4504 - accuracy: 0.5496
Epoch 6/10
9/9 [=====] - 0s 2ms/step - loss: 0.4504 - accuracy: 0.5496
Epoch 7/10
9/9 [=====] - 0s 2ms/step - loss: 0.4504 - accuracy: 0.5496
Epoch 8/10
9/9 [=====] - 0s 2ms/step - loss: 0.4504 - accuracy: 0.5496
Epoch 9/10
9/9 [=====] - 0s 2ms/step - loss: 0.4504 - accuracy: 0.5496
Epoch 10/10
9/9 [=====] - 0s 2ms/step - loss: 0.4504 - accuracy: 0.5496
```

```
Out[43]: <keras.callbacks.History at 0x28ac1f2d900>
```

```
In [44]: model3.evaluate(X_test, y_test)
```

```
2/2 [=====] - 0s 4ms/step - loss: 0.4754 - accuracy: 0.5246
```

```
Out[44]: [0.4754098355770111, 0.5245901346206665]
```

In [45]: `model3.summary()`

Model: "sequential_3"

Layer (type)	Output Shape	Param #
=====		
dense_9 (Dense)	(None, 64)	896
dense_10 (Dense)	(None, 32)	2080
dense_11 (Dense)	(None, 16)	528
dense_12 (Dense)	(None, 8)	136
dense_13 (Dense)	(None, 1)	9
=====		
Total params: 3,649		
Trainable params: 3,649		
Non-trainable params: 0		

In []: